ActInf OrgStream 006.1 ~ d20 Governance: an experimental governance en

OrgStream | d20 Governance: an experimental governance environment | 1:30:21

Hello and welcome. This is Active Inference org stream 6.1. It's January 22, 2024. And we're here

with the D20 governance team and project. So I will let them all introduce themselves. This should be

a very fun and interactive session. So thank you all for joining and looking forward to your

presentation and the discussion. Cool. Yeah. Thank you, Daniel. So I'm Seth Hostin, and I am a community lead at the Meta Governance Project and one of the developers and researchers with the D20 governance project. I'm here with Basil and Janita. And I'll let them introduce themselves now, maybe

Hazel and then Janita. Hi, thank you for coming to the live stream. My name is Hazel. I'm affiliated. I'm kind of in a medical community environment atmosphere. I am political theorist. I'm a researcher with D20. Before I pass to Janita, I also want to highlight another of our team member who isn't here. Vel is her birthday. Happy birthday, Vel.

Hi, I'm Janita. I'm a software developer. I have been involved in the DAO ecosystem working on some projects around DAOs in the past. And this project is my main sort of interaction with the MetaGov community. But I'm generally pretty interested in all the work that they're doing. And

I'm a researcher in governance, et cetera. Cool. So with that, we're planning to do a presentation today. And then the presentation will also have some interactive demos for people who are joining us in the stream.

And yeah, so with that, we'll get started. I should also say before we get started that I want to thank Daniel for inviting us to come and speak at this. I've known Daniel for as long as I've been doing the community lead work at MetaGov.

And I've had a lot of really interesting discussions with him about cognitive security, stigma G, and governance in general.

So I'm less familiar with the kind of resources happening at the Active Inference Institute. So really glad that Daniel and their community was interested in having us come and join so that we can present our work and learn from the insights of this group of researchers.

So with that, let's go ahead and get into the presentation.

Great. So before starting, there is this Discord link. And actually, it's not in my clipboard. So I'm going to generate a new invite.

And then I'll share it in our chat during Zoom. And then Daniel can share it in any, like, whatever way is applicable.

And then, yeah, so this is going to be a link to our Discord server, where the bot for this project is hosted.

It's the D20 Governance server.

And we're going to do some demos today. And so if you want to participate, this is the place to join for that.

If you want to join, come in and say hi. And yeah. So with that, I'm going to pass it over to Hazel, who's going to give us a bit of an overview for today's presentation.

Yeah. So the flow of this presentation today will be some demonstration of the functionalities we built, sprinkled with some conceptual thinking.

So we're going to start with giving you guys a background on where this, how this project came to be.

And we will introduce our first functionality. That's D20 Agro, where we're going to be interacting.

And we will demo the culture modules and talk a little bit about community rule, which is where and how this would inspire this project.

And we'll talk about two more functionalities. That's value, group voice, actually three more and continuous input.

And then we'll end with a Q&A and discussion.

So moving on to the next slide, a really brief overview of this project is D20 is a Discord bot that allows communities to play governance games in an AI large language model mediated environment.

So groups of people can come together to embark on a governance, what we call governance quest and experience a diversity of governance structures, decision making processes, power distributions and cultural dynamics.

So we're building the Discord bot, D20 aims to help lay the collective decision making foundations necessary for online communities to conduct their own governance experiments, do governance education, do future scenario planning and make real life decisions while bonding communities through play and a sense of playfulness.

So I'll pass it to Janita.

I'll pass it to Janita.

Janita

Janita

Yeah.

So actually, if you don't mind going to the next slide, I think there's a little GIF I wanted to reference.

Janita

Janita

Again, we'll see this is kind of the ail

and have it just work for sort of your community's governance.

But we also, at the same time, wanted

to kind of have this notion of gameplay

or allow people to experiment with these modules

in a very low stakes manner that you could then move forward

and take into higher stakes situations,

like actually using these modules to govern your community.

So the idea of low stakes is kind of where the governance game idea came from.

So we started out developing a more fully fledged kind of game experience.

And we were trying to design a narrative that took people through a series of different decisions that they had to make as a community.

And we're trying to be pretty honest to the original structure of community role

with respect to these four different types of modules.

But something we realized as we went along was that game design

and making a really kind of interesting and exciting gameplay experience

is a hard problem in and of itself.

And we were kind of trying to accomplish a lot of things at once, we realized.

So we decided to pare back the game structure a bit and sort of focus it around what we call

like these mini quests or governance experiences where people in the community can kind of like still experiment,

but in a sort of like on a shorter timeframe and with less narrative overhead,

they can still experiment with these governance modules.

So we'll get a bit more specific about the actual structure of the game later on,

but that was kind of the origin story of the D20 bot.

One other thing I'll mention is that as we were like developing this and doing play tests with different communities,

we got a lot of feedback around the sort of this LLM mediated communication feature.

So one of our original ideas was to implement these like implement the idea of culture models,

but in a more wacky manner by using LLMs to modulate how people are communicating.

So, you know, for example, we have like an eloquence types in the discord is translated into Shakespearean English.

And we found that like people were finding a lot of delight in these culture modules.

And so we started thinking more about, you know, how this game can allow people to experiment with communication

and LLM mediated information landscape or like how even the frustration that you experience when trying to communicate amongst sort of this like chaos of information can be interesting or can can teach us something about governance.

Yeah, next slide, please.

So actually, I'm going to just go back to the other side really quick because there's something I wanted to add on to what you were just saying there, Janina.

Something that's also kind of interesting about, you know, modeling the kind of communication system based around community rule is that the kind of implementation of the different kinds of culture modules.

So, you know, the kind of communication can vary quite a bit.

So if you look at this actually, well, you can't see it because it's not on culture right now in the GIF, but in culture, there's a values module.

And we'll talk about that a little more.

But one of the interesting things is that that's a quite broad category for structuring communication as like mediated and filtered through the language model.

So some of the modules are very narrowly defined and some of them are more abstract.

Another thing that's sort of interesting is that this tool was kind of originally designed to basically help communities define the way that their governance actually is rather than necessarily the way that they want it to be.

It's more of a retrospective tool rather than a prospective tool.

And part of the reason that is that it's like that is because the community rule itself is not actually programmatic in the sense that setting up these structures doesn't do anything.

It's not connected to any APIs or the web app doesn't even do anything other than just represent.

And so kind of thinking through how to translate this into something that's programmatic was actually a really interesting exercise.

And it's worth saying that too, because one of the things that kind of was also part of the impetus for this work was some of the past user research that we had done with communities who were interested in incorporating more governance tooling into their suite of online tools.

We, Val and I in particular did a lot of user research with community participants and leaders while we were doing some work on both MediGov's Gateway project and also PolicyKit and now Collective Voice projects.

And one of the things that came up a lot in those conversations was the need for some kind of primer or some kind of accessible way to even develop the kind of governance muscles or impulses that an online community would need to really make use of the tools.

And community rule is actually, at least in my opinion, a really great tool for introducing communities to what a very simple kind of representation for what governance could look like.

The drag and drop web app and the kind of simplicity of seeing your structure represented, even though you can actually get these very strange kind of tools.

very strange kinds of recursive structures.

But you could have values inside of a board, inside of a values inside of a board, inside of values inside of a board, just kind of infinitely recursing.

But it's a really accessible and approachable way of actually seeing what it means to have a structure, to make something that is often implicit.

And so that's part of the reason that we were taking community rule as a foundation.

And then we also heard that there is this desire to something that is more explicitly interactive.

A lot of tooling can sometimes be set up a poll and then let it run.

And you don't really get any kind of interactive feedback as a result of participating in that process.

It can feel like a lot of online interactions where you do something, maybe you see like a react number go up or down.

But largely speaking, you do something and you kind of don't even know if what you did had any impact.

And so this idea of kind of creating a way of introducing people to online governance through play seemed to be something that was coming up in our research.

research.

Janina talked about this a little bit, this sort of in terms of the kind of conceptual development of the project, moving from a kind of general governance tool and even a game to thinking of it more as a kind of surface for governance experiences.

And I think this is maybe important to linger on just a little bit because there's something about having a good experience with governance or a meaningful experience with governance that can really change your whole relationship to why you would even care about the topic.

And so even though, like, as we'll kind of get into our game is our kind of governance experiences are designed in some ways to kind of be, to reduce confusion.

They're also designed to kind of stimulate the imagination for what a kind of governance could look like and still have some fun in the way in order to.

To basically translate the experience of playing into thinking about implementing these kinds of systems in your communities.

The other thing is that for, I mean, one of the hard things about designing governance scenarios, games, environments, et cetera, is determining what the resources are that are actually being governed at all.

And it's very nice to kind of, it's very nice to kind of imagine this very abstract governance scenario.

But if there are no stakes in terms of resources, then there's really, it's not entirely true, but it's much harder to kind of devise a governance interaction space.

And so, you know, I think we started and maybe we'll pick up a little bit on this and some of the kind of prompting discussion later on.

But we started with this idea of kind of living in a kind of like post communication collapse type environment where people, for whatever reason, still needed to communicate over the internet.

But the ability to do that had degraded so much that communication itself was this kind of clarity of meaning was a scarce resource of sorts.

And the other thing that we were interested in was thinking about how to kind of like produce that kind of scarcity or that kind of fleeting sense of meaning.

And the way that we use cultural modules was one way of kind of artificially placing people in these kind of abstract communication environments.

So, yeah, I'll go to the next slide.

Janita will talk a little bit about the D20 Agora mode.

Yeah, so what you see going on here on the right is the D20 Agora.

So the game is basically broken up into two modes.

There's the Agora, or sorry, the bot rather is broken up into two modes.

The Agora is what we're looking at.

So the other thing is that it's kind of this asynchronous environment where all of the, you know, culture decision, et cetera modules are available to experiment with.

But there's no kind of active game or narrative going on that requires people's synchronous attention.

So we intended the Agora to be an experimentation zone and the jump off point for the,

the actual games experiences, whatever you want to call them.

And we'll talk about that second piece in a bit.

This is basically just a kind of GIF, but actually, because we're doing the stream this way,

it's basically redundant.

So just kind of to the next slide.

So we're going to go into a little bit of more detail about the way the culture modules are implemented and designed.

We'll start with Hazel and then I'll pick up a little bit after.

Yeah, I think culture module is one of the really interesting functionalities that we build and it's the way through which we mediate kind of all messages and communication through the switching on and off of culture modules.

So to give an introduction, I think culture modules is one of our most notable features and it works to influence the group's communication norms with what we call.

We picked up the term culture modules from the community rule where basically we use large language models to modify the content and the input of users' messages and temporarily place the users in the artificially constructed culture communication environment.

And a lot of these culture modules have pre-generated large language model prompts that the AI use to transform the text input of the user's message on Discord and give an output.

And we see a lot of the transformations happening here.

So if you type something, your language will be mediated and changed by a bot.

And with reference to whatever culture module that's active.

So these are the six culture modules that we've built so far.

And to just, yeah, we have Amplify, we have Eloquence, Obscurity, Ritual, Values, and Wildcard.

Amplify just intensifies the message sentiment.

Like if I say that, you know, today is fine, it's going to amplify it to be something like today is great or today is awesome.

The best day ever.

Eloquence makes messages more persuasive.

And in our formulation, they kind of turn the message into like a Shakespearean like poetry kind of way, which is really, really interesting to read out loud.

In the environment and obscurity kind of makes the message more obscure by like a scrambling of words.

But we actually found out like really surprisingly that it's still kind of readable if you stare at it for a little while, but it just makes it a little bit harder to read.

And ritual is interesting because it actually changes the meaning of your input.

It will align if you turn on the ritual module, it will like align your message and to like with the previous message that's being sent.

That if the previous person say like, you know, I love like tea or something and you say I hate tea, it changes the like I like tea to have this like really interesting mediated group consensus that's very constructed.

And we have values, we have a list of values defined in the orgware channel, you can propose revision to it, and you can check values with the check values function.

And the last one is a little bit special, the wild card module.

It transforms messages into the voice of the culture that's being defined by the previous group.

So this is like a special one we built in.

It's designed to construct like a new culture module every time a group successfully completes a build a community quest.

So the transformation prompts constructed using the text content of the decisions made through governance mechanisms made during the quest and made by the group.

And the questions the players are tasked with making a decision are related to the voice that's going to formulate and inform the what the wild wild card culture module is.

And we will talk a little bit more about the build community quests later.

And I'll pass to Sen to talk a little bit about the conceptual thinking behind using AI generated prompts.

Yeah.

I talked about it actually a little bit already, kind of maybe before it was scheduled in our presentation.

But I think one thing that we need to kind of come at it from a different perspective is, you know, we had people at DwebCamp when we did an in-person play test of this.

And they, when the eloquence mode was turned on, they ended up just like reading these posts out loud.

And afterwards, we had people sort of talking about how there was this, you can imagine a scenario where they're in a discord or some kind of communication platform.

And people are basically talking past each other.

They're on their own individual tracks.

And there's the idea of kind of coordinating or responding to what the other person was saying is basically impossible.

And they said it would be interesting if they had the ability to go into a kind of eloquence mode or step into a channel where all of the communication was structured in this way.

And I think this kind of points to the, something that I think we want to come back to in the discussion, but this kind of question of producing artificial stability in the environment.

And like, I think often it's maybe speaking for myself, I think often when I'm posting, I become very attached to what it is that I'm posting for a couple of reasons.

The content is, I can really only express myself through text usually.

And so a lot of the nuances, of course, of how I might actually express something are lost in the communication platform or the medium.

But also there's the kind of sense that what I am saying is also kind of part of a long archival history and part of the kind of reputation of the media.

The kind of the kind of reputation that I've either explicitly or implicitly developed within a community.

Well, the chat and being able to kind of create a situation where everyone who's posting is kind of conditioned by the same logic of communication.

I think helps to serve a coordinating function, even if it's a frustrating coordinating function, but it also serves to give a little bit of distance between what it is that you're saying and what it is that you're communicating.

And then how you continue to communicate as a result of those transformations.

People, including myself, also notice that the way that you type changes because you're in some ways trying to kind of imagine how will your text be transformed by the language model.

And I think this kind of introduction of a kind of flexibility in this kind of artificially stable system allows for a slightly different approach to thinking about how you make decisions with other people online.

I think it increases the kind of complexity of making decisions because the ability to accurately communicate is made more challenging, but it opens you up to new ways of relating to your peers in that environment.

And so I think one thing that we're sort of interested in is how do these culture modules kind of influence the decision making patterns within online communities.

So at this point we were planning on doing a demo, but the discord is already quite active and the culture modules have been pretty thoroughly engaged with.

But one thing that we were interested in is sort of right now, the discord bot and the project is hosted on our discord server.

We still need to go through the process of setting it up as a bot that discord servers can install.

If there's interest, the code is open source and you can set the environment variables and actually just put it in your own discord.

But the reason that we're saying this is because like right now it's hosted here and this is a kind of decontextualized space on the internet in some ways.

Like people from your community are here right now and you're missing all of the context of the channels, the kind of flare that you might have set up with your account, the ability to interface with the kind of specific environment that you operate in.

And we're actually really curious to see what that is like because right now it feels kind of distanced.

And also there's no real incentive or magnetism to this particular discord server.

The only reason you would come here is to like play with this, but it's not integrated into the place where you are in community with others.

And so to kind of stimulate or simulate that experience a little bit, we thought it would be interesting to have the alongside the rest of the presentation.

People having a conversation about something to do with active inference.

We were thinking that it would be interesting to have just a kind of general account from the active inference community of what it is that they're interested in with active inference.

And how they understand the term.

And then to maybe kind of just carry on attempting to have a conversation in the way that you might normally do it in your own discord.

Just to kind of give a little bit of that sense of the kind of transformations happening in a contextually meaningful environment.

So I'm not going to pause on this, we'll just go to the next slide.

And this is going to be Janita talking about one of the specific culture modules, which is values.

So did you intend for there to be a few moments for discussion?

For discussion, like in the discord?

Yeah.

I'll read a question from the chat.

And also we can talk a little bit about active inference or continue on.

But I'll read a question first.

Upcycle Club writes, can you combine two or more modules into a new fused one?

Yeah.

I think that's actually what we were experiencing a few minutes ago.

So you can have any number of culture modules running at the same time.

Right now I think there aren't any culture modules running.

But if, for example, we wanted to turn on amplify as well as eloquence.

Anything anyone put going to be run through both of these prompts.

So the sentiment of your message will both be amplified and transformed into Shakespeare in English.

One of the interesting things to point out from a technical perspective, just how the code is actually implemented is the order in which the culture modules are activated has an impact on the way that the text is transformed.

It's the culture modules are stored in a list and that list is being read through in order.

And the order of the list determines the way, the sequence of the processing.

So you have your input text.

Let's say we have amplify then eloquence.

So your text goes to the amplify transformation and then the output of that goes to the eloquence transformation.

So there's even kind of an ability to create basically amplified followed by eloquence and eloquence followed by amplify are two kind of different functions.

And I mean, that's like in a technical sense, but there's another way in which, in my experience anyways, the eloquence module kind of has this way of overriding more or less all the way.

More or less all the other modules.

It's so distinctive in terms of the way that it's characterized that I often find that if I want to get a sense of the flavor for the other modules, I have to deactivate the eloquence module.

There's also, you know, I can't remember, but I think in the, the Agora, you're not even the ways of changing the culture modules isn't exposed or maybe it's not exposed in the quests.

But to answer, to also kind of answer what I think the person in the chat might have been asking is there's no way to kind of make like to specify a module that isn't already specified.

So you're kind of, you kind of have what you're given.

The wildcard module changes as a result of the quest.

And it's only temporary.

So after another group plays the quest, then the wildcard module changes.

We would like to explore the idea of more or less being able to kind of specify the way in which the large language model is transforming the text.

The text that that I think there's a lot of ways that the kind of system was architected.

It was very kind of organic and not predetermined.

Maybe at some point, you need to be interested in talking about that.

But the way that it was developed didn't kind of have the path dependencies that led us to a system where you could just create your own module on the fly.

Awesome.

And just one more question from Pablo FM.

He wrote, if in our role playing game, we have our set of values and our lingo and lore, can it be adapted?

So one thing that we kind of mocked up a very simple like web app that would allow people to kind of basically define their own quests.

And it never really is finished being developed.

But we're interested in definitely being able to let people construct their own role playing scenarios, define their values in advance.

Maybe you can have them stored in a database.

Every time that the server where the bot is hosted is restarted, all of the data is wiped because it's just stored defamation of the way that governance is conducted in that community.

So you can imagine joining a discord server and then in order to participate in some collection of channels, you first need to go through a governance quest.

And part of that would involve communities being able to define their governance in advance in the way that their quest is structured so that when people join, they have a stable reproducible experience.

And this is also the idea of kind of working with community rule.

We explored this idea of like actually exporting the,

when you play a quest, you have these different opportunities to change which culture module you're interfacing with,

what decision-making method you're using.

And we had this idea of kind of showing people how they,

how their governance stack or their governance structure changed over time,

and then allow them to basically export the, any slice of that,

and then use it as a way of generating a new quest.

So these are all like kind of future development ideas.

But the short answer to Pablo's question is no.

Awesome. Continue on.

This is really fun. Really fun.

Definitely bringing a smile to many people.

And there's a lot of, a lot of seeds that you planted, like the context lists,

the decontextualized or the recontextualizing of the dialogue in a space that's new to us.

It's in your lobby and there's all these different formal informal structures.

So we'll return to some of this, but let's see a little more.

Awesome. Cool.

Okay. Yes. So Janita to talk about values.

Yeah. So,

so another one of the culture modules we have available is called the values module.

So I think Hazel or Aracent have touched on this earlier in the presentation.

It's basically a way that you can sort of softly enforce the values of your community.

So anyone can propose.

So there's kind of a fixed number of values that we have.

And anyone can propose a revision to any one of these values to replace it with a new value.

And then what you can do is when someone posts a message,

you can see if that message is in alignment with the values that the community have chosen by running this command called check values.

So I'll just demonstrate that on a message.

So we see at the,

I think someone ran the values command.

So the community defined values we have are inclusivity, collaboration, trust, no respect and intentionality.

So I'm going to go ahead and run check values on upcycle club's message.

And the values analysis appears to be that this message does indeed align with our values.

And then there's a description for why the bot thinks that it aligns with the values.

Yeah.

So just another building block that we've included to sort of allow people to experiment with like how things can be made more.

things that are generally implicit can be made more explicit.

And,

okay.

Yeah, I guess we can try if,

Daniel, if you want to try the proposed value revision command,

just propose underscore value underscore revision,

and then it'll take you through the flow of replacing the no respect value.

Maybe go,

go through it on your side,

just cause that's what screens captured,

but this is epic.

Yeah.

It's very flexible.

And I'll just read one comment just while you're setting it up.

I'll read one comment from Dave.

He wrote,

here's a very practical application of the kind of transformations you're demoing.

Please consider working on this at some point,

glossing a discourse that's expressed in one vocabulary into a discourse expressed in a different analogous vocabulary.

Identifying corresponding terms and concepts in two fields or in two approaches to a given field is expensive.

Deploying the result of such discovery will be useful on its own and a spur to the effort to develop more correspondences.

This would involve a lot of under the hood activities required by full machine translation, e.g. German to French, but could be driven by much simpler rules, aka patterns.

So, thank you, Dave.

All right.

So you're typing up respect.

And justifying it.

Yeah.

So, so you'll see that we all have the, the option to object.

It's kind of like the past by default, unless there are any objections to the values within a certain time period.

Um, so provided that no one, yeah.

Lazy consensus is the term for that, uh, decision-making process.

Um, provided that no one objects to this value, it'll be, it'll be added to our list.

Cool.

How long is the time?

Like you could set different amounts of time.

Yeah.

Um, yeah, I think the default is like a minute.

A minute.

Yeah.

Yeah.

You could set the time, like you could program that, but, uh, we haven't partially because one, the, someone could set it for like 60 minutes.

Um, or someone could set it for like one second.

I mean, you could like specify the, like the boundedness of what's allowed.

Um, but usually because anyone who objects to one of these values is going to be able to pretty quickly realize they object.

Um, our sense was that you really don't need that much time, uh, for this decision.

Yeah.

Oh, it totally makes sense.

It just reflects how that could be a one month review process, like in a DSI scientific review setting with the right interfaces, or it could be really rapid, like feedback from the individuals and their time bound activity.

And then that feeds back into the game, like new rule.

Here's how we propose new rules, or here are the options for rules, or here's how we're going to apply which filter.

And that just adds so many plot twists even before active inference or without any kind of cognitive modeling at all, or systems modeling, just approaching it from the pure gameplay.

It's really fun.

It's really fun.

And people in the chat are like, yeah, already pointing out a lot of ways it can be used to play, if not be more useful.

Yeah.

One, one thing to point out that's kind of hidden at the moment, uh, is, so let's go to this value analysis of upcycle club.

Um, uh, where are they?

So they were, um, aligned.

And if you go to their, uh, profile, um, uh, no, um, there should be something that, uh,

I saw that there was a misaligned tag on Friendly's profile.

Uh, yeah, yeah.

Uh, I don't know if it's doing it right now, but, um, I think we might have taken it out at some point because it wasn't part of the full feature set that we wanted to ship.

Um, but, um, this function can assign an aligned or misaligned, uh, tag based on the results of the analysis.

And then you could do, do whatever you want with that.

Um, from a programmatic perspective, one idea would be that like, if, uh, you as a community are aligned, um, any user that you're aligned with, um, you are either,

uh, not affected by the disruptions in language that are caused by the culture modules, um, or the effects are less intense.

And so there's an incentive to, uh, remain aligned if you want to be able to effectively coordinate with your peers.

So, um, the, the demo thing, the demo idea that we were going to, um, explore was, um, asking, uh, people in the chat to answer the question of what do you value in communication environments?

Um, perhaps from the perspective of active inference.

And then we can do some, uh, value, uh, analysis, um, even propose new values.

Um, and we can sort of see this kind of, uh, iterative feedback layout.

Um, so maybe we'll just leave that again, like, just keep playing with that in the background in the chat.

Um, it feels, uh, you know, something about just stopping and watching the chat and not saying anything feels, uh, too strange.

Um, so, um, uh, yeah, we'll go to the, um, building a community mode.

Um, I know that I was just talking about this, but I feel like I've been talking for a really long time, even though we like structured the speaking setup.

Um, Hazel, did you want to talk about this at all?

Or do you want me to still pick it up?

Go for it.

Okay.

Um, yeah.

So, uh, building community, uh, we have a build a community quest.

Um, it's, uh, we actually explored a bunch of different quests, this kind of like post communication collapse quest.

Uh, we explored one that might still be in the repo somewhere called the Josh game, um, which was a, uh, kind of, uh, actually, Hazel, can you talk about that just very briefly?

Like what the Josh game was?

Yeah, um, I think the Josh game is like, um, it's like intending to be an, uh, a game that's placed on individual identity.

It's kind of inspired by, I think one of us saw like a news online about the like IRL Josh game that a bunch of people are named Josh.

Um, and they kind of had like a, just like a fake like noodle sword fight, um, for the name Josh, where the last person standing who like won, like via rock, paper, scissors is the last person.

Who's crowned Josh.

Um, and the rest of the people like lost their ability or lost their right to be Josh.

Um, and I think like through, we kind of made some version of that where we all enter the game with our like avatar and name replaced with like a random Josh.

It can be like Jocelyn Josh or like Joshua Tree Josh or whatever else.

Um, and, and then we kind of like compete.

We kind of give spiels and, um, speech in this like kind of word-based contest where we kind of compete to be who's the most worthy Josh or who's the last Josh.

Um, and yeah, it's just like, um, like a playful way to experiment with who's who, um, how do you, like what's in the name and how do you like kind of like establish or spin out like a little playful identity.

And we imagine like a lot of different ways we can potentially play with that.

Um, but, but that's the, that's the general, that's the general like, um, playful context of the Josh game.

Cool.

Thank you.

Hizzo.

Uh, another funny, uh, thing about that game is, um, uh, Joshua Tan, um, uh, was the executive director of MediGov at that time and, uh, is still currently at MediGov as a resource director.

Um, and, uh, we actually never played this with him, but it's, uh, his, uh, GitHub handle is the last Josh.

Um, and so, I dunno, there's something very funny about spending one's time, uh, making a game that like is not even directly referencing someone, but is yet nevertheless connected somehow like organizationally.

So a little bit of, you know, meta for, for the extreme posterity.

Um, so to, to the build a community quest, uh, yeah, 25 minute narrative led game.

Um, and, um, you can prompt it in the chat.

Um, I mean, you can try it, uh, it's been a little while since we've actually done it.

Um, the bot might break.

Um, so, and also it takes a long time to do, but, um, basically, um, you can set up the parameters of a quest.

Um, you can decide how many players need to join, um, and whether or not images are generated as part of the quest.

And then you go and you play.

And the goal in the, in the game, the quest is to basically kind of, as a group collectively establish, um, yourself as a new group.

To give yourself an identity of some sorts.

Um, you have to answer questions like what's your name?

Uh, what's the main topic of interest that your group talks about?

And in what way does your group can speak or communicate to each other?

And for each of these questions, there are proposal functions.
Um, you can revise proposals.
Um, there are voting, uh, decision-making processes.
There are some deliberation questions that are sent.
Um, and there are also different methods for making decision.
Um, at the moment I think we've implemented, uh, a couple of different methods.
uh two or three like a very small amount of decision making methods um and then there's a
fourth kind of implicit feudalism uh method um which is if the google players um are unable to
make a decision after three tries the bot basically just picks a decision at random um
um partially because um i don't know it's it's hard to know how deep we want to go into the implicit
feudalism thing um but basically one kind of like background thing that's also connected with all
this is the idea of um thinking about the ways in which the platforms we use to communicate with each
other are uh structured and nathan schneider has written a really nice piece called um the implicit
feudalism of uh social media and it was looking at like the history of the server architecture
the server client architecture um from as far back as the vbs era up to now and his argument which i
won't go into a lot of detail about is basically this architecture and the ways that we communicate
in general are basically uh representative of a kind of uh feudalism um it's not explicit but it's just
implicit in the way that we communicate and um one of the theses is that we could give communities online
a uh different ways of um governing or setting up governance systems that are not inherently feudalistic
and this is kind of where the modular governance and the meta governance angle of this project comes
into play um and a lot of the the game is about kind of introducing uh people to the idea that they
could change uh on the fly how they make decisions um they can have agency over um the way um their
communication happens in a way that is not uh is circumscribed by the developers of platforms and to
be honest it's also circumscribed by us um it's very difficult to create a system that doesn't have any
kind of predetermined uh structures in place um but for example in the game if if you fail to make a
decision about any of these topics then you're given the ability to uh change the way you make decisions
um and then if none of that works then a decision is just made for you and this is in some ways meant
to kind of bring to the surface for players the ways in which um they can and cannot exercise agency
um when communicating and the other thing to say about this this question why you go through this
process is that um all of the answers to these questions are then fed into a template for an llm
transformation module this is the wild card module and what that means is that like one thing that i think
is there there are a lot of experiments with kind of uh multi-user input uh interfaces into llms but not so
many um and i think one thing that's sort of interesting to think about is how can we open up um an imaginary
um and i think that's a little bit about um i guess two things one is um having more uh collective uh inputs
into um the kind of uh way that the large language model is interfacing with uh communications um but also
kind of um being clear through experience the ways in which the way that we communicate and the
decisions that we make have an impact on the ways in which our environments uh change around us
so this is kind of maybe like the stigma g kind of thing a little bit like you have these like uh

interactions these like uh uh pheromones that people are leaving behind which are their decisions
and then you kind of experience the result of that in the agora and that also changes the way that you
interact in the agora and maybe the way that you interact in the build a community quest
um so i mentioned that if a failure if a decision fails um then you have the ability to change it
and the way that that happens is through a kind of continuous input mechanism um
um it there's a kind of multi-user uh input uh uh there's a way of uh inputting your uh preferences
for a decision um and then once um a certain level of preference collectively has been expressed then
uh a decision is made and the mechanism is triggered um it's also a display of community conviction
uh it also just kind of acknowledges the fact that there are governance failures i think one thing that
kind of comes up sometimes and discourses around governance and also the perceptions around governance
is that governance is kind of uh when we talk about it as a system um it can sometimes uh
uh or a model it can sometimes be perceived as this like static set and fix and don't think about again
uh experience and sort of a lot of the literature that kind of thinks about especially the emergence of
meta governance and governance in general as a kind of development of government um is basically like
hyper focused on the ways in which failures manifest and and decision making um so we wanted to make sure
that when this happened um there was a way of exercising agency and and kind of giving players the experience of
steering or uh responding to the environment another thing to maybe say just quickly about this is that
this kind of uh mechanism is also one of the like most uh popular uh often cited uh kind of uh ways of uh
references for this is twitch plays pokemon um which is basically a emulator of pokemon uh uh set up on
uh on twitch and then some python binding so that anytime someone types uh up in the chat then uh an up uh command is uh sent to the emulator and the character moves up um and so you can have
a chat that's basically just collectively distributedly controlling the game and uh this is cool
um and extremely chaotic um but there are maybe like two scenarios in twitch place pokemon or like in pokemon
where you really have to press the same button maybe like six times in a row in order to even move to the
next part of the map um and so like you can't stay in the kind of pure chaos mode um or else you'll never
complete the game and so the developer of the python binding uh set up this way for people to oscillate or move
between uh what was it like uh democracy and anarchy um and when you're in anarchy mode it's exactly like
we just described it's all continuous just taking the input um and then when you're in democracy mode
there's a kind of time buffer in which um all of the responses are collected and whichever response
gets the most input in that period of time is the response that's triggered so it sort of slows down the
the reading of input um and allows for groups to coordinate and actually get past these situations
uh where they couldn't otherwise and then it's a super interesting project it's very old at this
point as well but it's highly recommend people who are not familiar with that already to go to look at it
another great reference in this area is uh trusts uh moving castles project which is also looking at

community controlled games through twitch and is a kind of more contemporary example um
yeah so we wanted to demo this and you can see that it's already happening in the discord chat a
little bit um people are typing a culture module and then plus or minus one and then once it passes a
certain threshold then uh that module is activated so we didn't have any prompt for this one uh so maybe
we'll just kind of move to the next section unless anyone wants to pause and discuss something since
that seems to be a pattern here demo discussion
i think one quick thing i would add is we told we built the cool down time because i think when we
were playing it we ourselves find ways to break it and cheat on the system um i think that's also part of
the fun of building the game of thinking about not not just thinking about ways that we can play but
also to think about to play through the disruption by disrupting the disruption um which was really
interesting which is why we built the cool down so nobody can spam and just get what they want at least not immediately
yeah that's a great point uh yeah yeah thank you um anything you want to add maybe janita or
okay um so progress um we've done a playthrough of this at demo camp we talked about that a little bit
um we did a kind of annotated uh play through uh like an annotated version of that playthrough um i'm
not going to show it because it's not set up um i can't remember where the link is actually but
anyways um that was a really interesting process and we should make a slide for that next time um we also
did a presentation as part of a symposium in geneva called uh the sea shore of endless worlds um and we've
also presented this at a uh one of the weekly meta governance seminars so um and now we're here um which
is very exciting um and really looking forward to this stream existing because i think this is uh
a very fun uh engagement with the project um so yeah with that um with that maybe we would move into a
little bit of discussion even though we've had this nice kind of uh interjection discussion so far um
and we were thinking of um like in what ways is this kind of uh project um how does the active inference
institute community kind of see relevance in it um i guess in terms of both maybe setting up experiments
or um kind of um kind of using it as a complement to or the subject of research and then another thing
that we're interested in is kind of thinking about this kind of the thing that we alluded to earlier this
kind of like uh communication collapse and like a kind of post llm or whatever uh kind of uh technology
were posed at this point uh communication environment where um the like there's this idea as far as we
understand we're not experts in active inference that like one of the kind of reasons to study and model
this um in this field is to kind of be able to uh anticipate and respond to unpredictable environments
somehow like the idea of active inference and i'm curious if this is correct is to kind of produce a kind
of more stable system than uh we might um otherwise uh design and and this game in many ways kind of
purposely introduces um uh instability into the communication but it's also there's there's
those disruptions are themselves kind of meta stable environments and so there's this kind of
strange interplay between stability and instability and the ways in which this kind of experience of
of um expanding the kind of governance imaginary for how we might interface in a more ambiguous
uh uh systemic uh uh systemic uh context so um we'd love to discuss that a little bit and then to maybe um

leave

uh everyone with some information about us um here are our names and our respective roles in the project um

some of our contact information um and then some links to our projects and um we should also mention

that the project was created and uh supported by the meta governance project so yeah that's our

presentation thank you so much for having us and um i'll leave the discussion prompts up and uh we'd love to

see where the conversation goes awesome well thank you for presenting this

this like there's a lot to say a lot of places to jump in it feels so natural despite the tech in some

ways being so new in practice and in our hands but even though you all have made this it feels like

people just slotted into it and i think it it does touch on a lot of active inference ideas like on a more

more technical but also on a more informal level like just starting with the informal the idea that

there's multiple agents or multiple entities with varying blends new mixtures new filters and overlays

of bodied and digital all these different modifications and and layerings and identities

so this kind of like open cyber complex identity engagement space and the way that still there are these features of

ecological psychology and that those fundamental features like relating to sense making on the inbound

and decision selection on the outbound we can understand whether it's like leaving an emoji or the words

that are produced all all of that is behavior and it's behavior um that's different but with our body

so there's all these kind of like rich dynamics so the field provides more questions than answers on those

fronts but shows like how far the gap is between the general unified approach and how deep the last um

the general um mile is in application but this really like connects it from what otherwise could be super

abstract about like multi-agent online interactions which is like what the research papers in medigov or in

active inference how they'll describe it and this like really connects it and it's like in this discord

so it kind of becomes embodied in the real chat so that was one thought another thought was like

it it shows that words matter and that our speech acts online matter by representing them in this kind of fun

and unique way there's other ways to represent people's speech online like going back and curating older posts

that was something we had done or seen in the medigov slack and then here with some new shiny affordances

from the language models and all of that um to have words that have informal import with the sentiment that

we convey uh filtered or not and then also to bring in acts that have kind of quasi-formal import

like actually changing the rules of a game at some level of analysis and and being able to hold the fun

and playful and comical and that whole aspect and also seeing that that's even paradoxically or not

like what's important for high reliability distributed settings or for establishing trust or

new ways of communicating when there's change

thank you for yeah thank you for letting us speak i think to be fully transparent i think we're still very actively

uh this is not a finished project and we don't uh present it as a finished project i think we're still

very much in the process of imagining um use cases not just use cases but places where um this can be
like immediately useful or playful or make sense right because i think today we're really fortunate to
demonstrate it in a way that feels that has the split screen that demonstrate in a way that feels very
calm but we've also had like playthroughs and demonstrations where when it needs to be like put into
like say like 10 minutes it's really confusing it's just like a bunch of messages blowing up and
everybody's speaking um and i think yeah some users think it's really kind of chaotic um yeah i think
um yeah i think like we yeah we we really appreciate speaking today but would really still love to hear from
you guys um what are some of the use cases um and even like modules that you find that's like most
accessible um or any research potential that could potentially um like be interesting stemming from this
or other ways that we can i think play together like for the purpose of like governance spirit spiritedness
cool cool if someone writes something in the chat we can read it but just one thought that's kind of a general
approach in active inference modeling which is analogous to some of the work we did like in 2022 with dsci
distributed science decentralized science is take a look at systems that are often engineered from the top down
and then think about it from the bottom up in terms of the sensory inputs and then the capacity for action
the affordances of different agents in the system and so take that kind of view from the inside so one
kind of scaffolding of an approach that is pretty general cross active inference settings is like
kind of um kind of ail
so make a kind of scheme for who are the players and then with the players plus the board
on that's kind of the board game setup and it's a different board and different players
but i think even with the title of the project you kind of allude to that role playing
um yeah that's really interesting um yeah the title is also kind of uh a funny one um because
there are no dice involved um uh but it's it's catchy um the yeah i think having like more
persistent um consequences uh is an interesting uh avenue i mean the kind of aligned misaligned
uh role tagging that happened that we showed is kind of pointing a little bit in that direction
but um yeah like imagining like uh yeah kind of blank server the first thing you do is interact
with this and everything that you learn about the server as a result of the decisions and
interactions that you have um through this would be like a really interesting way of kind of
producing a somewhat like labyrinth and non-stable uh like constantly evolving uh environment for
people to interact in um and like yeah that's a great idea um it's really interesting
i mean that is newly on the table and so it's exciting to think about um digital environments
whether they're personalized or shared that have a defined construction to them so it's legible here's
how we're modifying the incoming posts or something so there could be situations where that's exposed
or like transparent or open source and then there could be other situations where for one reason or

another they don't and they would be accepting the trade-offs with or without sharing that kind of
information and it it brings like new um nuance to all the filters that we do receive information through
and as that becomes more post or multilingual then there's a there's so many new ways that this can
come into play because people are hashing these things out and they are having a model of who's
aligned and not and so making some of it visible can maybe that's a disconcerting look in the mirror
so not every representation is probably right for every setting but in like different settings there
will be people who love this mod facilitator meta governor role
and you're you're producing tools that help bring that to the discords where people really are
yeah uh there's there's two things that are interesting somehow there one is um i mean actually
working with the discord uh interfaces and also the api is actually kind of a nightmare
um like there's a there's a function that we have that um does a bunch of like text streaming
so it's taking a message and then like sending a couple words chunks at a time to kind of produce
this effect that the message is being written by some external agent um but there's a pretty aggressive
rate limit on how quickly you can edit messages um in discord um so there are some kind of
like constraints there um and i think the there's um there's something about this um i think i can
i one thing that i can sort of see is that there this this kind of model that we're talking about
of like a kind of really uh emergent and flexible and somewhat kind of unpredictable uh structure
for and and and a structure where no one person really has the same view on the system there's no
uh holistic view um maybe even as an admin um then like you're opening up so many kind of
unpredictable moderation questions i mean we already have enough challenges um or at least the discourse
will kind of point to the the kind of numerous moderation challenges that we have uh in kind of
maintaining uh our disc our discursive spaces aligned with some kind of sense of community value
um but like the ability to kind of have uh an environment where no one has the same view kind of
opens up this like huge like how do you even conceive of moderation in that way um there's there's also
the like uh there's also this um this problem of information overload um so we explored like
having the the generated response um actually show you the information um like the the original input
using like a spoiler tag um but it was just too just too much like it was visually cluttered um
and another thing that actually happens here is like i want to figure out how to like have it match
token length so that the input and the output are kind of like correlated to each other in length
right now there's this kind of uh it it will elaborate so much that it breaks the natural flow of
communication um and it introduces a kind of temporal friction and so like more moderation
uh to kind of uh to kind of expose what's happening in the background and it naturally puts a kind of like
break in some way on the kind of way in which we're communicating which could you know be beneficial or
uh or not thing and then yeah like this dystopian thing yeah i mean there are a lot of ways this could be
applied that are uh not benevolent um very like authoritarian um in nature so it's uh it is there is also this kind of
yeah i think the yeah the moderation surface just becomes so difficult to figure out how you actually approach
yeah well one interesting approach there is that now with these llms the space of moderation is in the

semantic not just the syntactic so a lot of discord bots or trending lists or all these things are based upon the syntax of like how something is spelled or actually like how word is used with all the kind of false positives and false negatives now we're in a different space or mode of false positives and false negatives with the semantics and that's a different and a more latent space and one that really gets at the cognitive diversity and different people's understanding of different words and decisions and events but that's kind of the real space as well and so making the maps of that cognitive semantic territory instead of saying okay no name calling here's the list of names now that is going to be elaborated on from something and and maybe uh and also just so new vistas for how people want to communicate and engage or not and like you brought up massive opportunities and challenges for attention and overload and just the pace of all this augmented text is extremely high

and so it's very intense also though people do a lot of intense digital experiences

yeah it's so true i think there's like uh some relationship between the intensity of digital experience that you're uh comfortable or have previously experienced and how quickly you pick up this this spot um we've done this in some other contexts where people are just like i have no idea what is happening um and i think if you are if you tend to be like extremely online um this feels very natural and if you're not it's it's somehow just way again it's uh it's too much um so that's also

it's also interesting from a governance perspective as well it's like i mean i can't make any real extrapolations but i wonder if in terms of the kind of the viability of online community-based self governance if there's a kind of correlation between the kind of like intensity of uh digital experiences that that the community is uh game for and how broadly effective they will be or adaptive to online governance they will be um whereas communities that are less active or less digitally intense um would generally find online community self-governance um harder to adopt

that's super interesting and communities that have variation amongst their participants in that as well

well this is an epic part one and kind of

play through i guess with this like maybe as we kind of close where do you want to go from here or what would be something cool in 2024 to get to or try um i don't know maybe we can each kind of try to give an answer to this um uh i think one thing that would be really cool is like

okay cleaning up the code base um so that it's like easier for other developers to to actually extend the work that we've been doing and contribute meaningfully um i and also getting this set up so that communities can actually experience it in their own contexts rather than just in this server so if there's anyone uh who is part of this community or who's watching who's interested in um kind of kind of sifting through the code base and figuring out how to contribute um that would be really cool

um yeah i think that's that's kind of my thought i don't know hazel or janina anything from you

yeah i i'm on a similar page um i really think of this bot and this tool right now as like

um yeah i think um yeah i think it would be great to figure out ways to extend this to make it more accessible for for people to continue to experiment and build upon what we've done

yeah i feel like i really wish that more people watching this live stream or not will just come talk to us and say hi and tell us how we can make this more useful or practical for you and that

would make us really happy and i think another thing that um a topic that i'm really interested
in is what kind of thing is randomness in governance um how does like being random work
and having a function that assigns like a random value work um i think there there are quite a bit
of we build like a kind of random overlord function um in build a com and in build a community quest
um but i feel like that's something i really want to explore more especially with regard to
um what's the relationship between um randomness and agency right like what's the setup that that
gives user like the agency the feeling of agency of governance um that's something i'm interested
in thinking through more um and hopefully with this project
cool super exciting i hope people in the institute discord and elsewhere
tune in to this kind of cool way to think about our shared services as governance services really
converting a space and taking everyone on this journey that's so
in the moment that it's kind of like intense and flies right by so it's been an epic stream
and maybe we can play around with it as you all continue to develop the framework so thank you all
till next time yeah thank you so much for having us cool bye
thank you